



# Distributed Computing and Introduction to High Performance Computing

Imad Kissami<sup>1</sup>

<sup>1</sup>Mohammed VI Polytechnic University, Benguerir, Morocco

May 10, 2023

## OUTLINE

- Why HPC?
- What's Supercomputer?
- Data locality
- How to make Python Faster
- Parallel models s
- Performance metrics

## OUTLINE

- The flood of Data
- Big data problem
- What's HPC ?
- Typical HPC workloads
- Data Analytics Process

## THE FLOOD OF DATA

### **In 2021**

- Internet user  $\sim 1.9$  GB per day
- Self driving car  $\sim 4$  TB per day
- Connected airplane  $\sim 5$  TB per day
- Smart factory  $\sim 1$  PB per day
- Cloud video providers  $\sim 750$  PB per day

## THE FLOOD OF DATA

### **A self-driving car**

- Radar  $\sim 10 - 100$  KB per second
- Sonar  $\sim 10 - 100$  KB per second
- GPS  $\sim 50$  KB per second
- Lidar  $\sim 10 - 70$  MB per second
- Cameras  $\sim 20 - 40$  MB per second
- 1 car  $\sim 5$  Exaflops per hour

# BIG DATA PROBLEM

Too much data




Not enough computer power, storage or infrastructure






# WHAT'S HPC?

Leveraging distributed compute resources to solve complex problems

- Terabytes → Petabytes → Zetabytes of data
- Results in minutes to hours instead of days or weeks



# TYPICAL HPC WORKLOADS



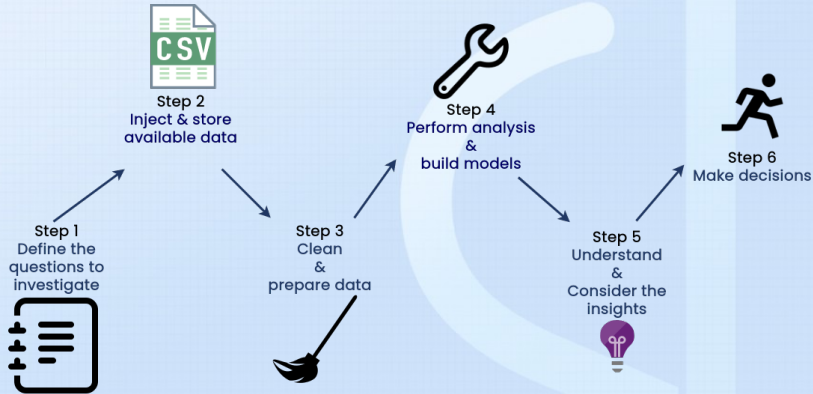
\* Source:

<https://www.xilinx.com/applications/data-center/high-performance-computing.html>



# DATA ANALYTICS PROCESS

Inspecting, cleaning, transforming and modeling ➡ decision-making.



## SUMMARY


- Larger datasets require distributed computing
- Several open source HPC frameworks available

## OUTLINE

- A brief introduction on hardware
- Modern supercomputers

# ❏ A BRIEF INTRODUCTION ON HARDWARE

## Modern architecture (CPU)



# A BRIEF INTRODUCTION ON HARDWARE

## Moore's Law

- Number of transistors: from 37.5 million(2000) to fifty billion(2022)
- Cpu speed: from 1.3GHz to 3.4GHz

### Moore's Law: The number of transistors on microchips doubles every two years

Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important for other aspects of technological progress in computing – such as processing speed or the price of computers.

Our World  
in Data

#### Transistor count

50,000,000,000

10,000,000,000

5,000,000,000

1,000,000,000

500,000,000

100,000,000

50,000,000

10,000,000

5,000,000

1,000,000

500,000

100,000

50,000

10,000

5,000

1,000

500

250

125

62.5

31.25

15.625

7.8125

3.90625

1.953125

0.9765625

0.48828125

Year in which the microchip was first introduced


Data source: Wikipedia ([wikipedia.org/wiki/Transistor\\_count](https://wikipedia.org/wiki/Transistor_count))

OurWorldInData.org – Research and data to make progress against the world's largest problems.

Licensed under CC-BY by the authors Hannah Ritchie and Max Roser.

## ❏ A BRIEF INTRODUCTION ON HARDWARE

### CPU vs RAM speeds



# ■ A BRIEF INTRODUCTION ON HARDWARE

## Common Processors

| Processor                                   | Launched | Nb. of Cores | Freq. (Ghz) |
|---|----------|--------------|-------------|
| Xeon Platinum 9282 (formerly Cascade Lake)  | 2019-Q2  | 28           | 2.6-3.8     |
| Xeon Platinum 8376H (formerly Cooper Lake)  | 2019-Q2  | 28           | 2.6-4.3     |
| i9-12900H (Mobile, 12th generation)         | 2022-Q1  | 4-16         | 3.8-5.0     |
| i9-12900KS ( Desktop, formerly Alder Lake ) | 2022-Q1  | 8-16         | 2.5-5.5     |

Table: Some Intel processors

| Processor                   | L3 cache | Nb. of Cores | Freq. (Ghz) |
|-----------------------------|----------|--------------|-------------|
| AMD EPYC 7773X              | 768 MB   | 64           | 2.2-3.5     |
| AMD EPYC 7763               | 256 MB   | 64           | 2.45-3.5    |
| AMD Ryzen 9 5950X (Desktop) | 72 MB    | 16-32        | 3.4-4.9     |
| AMD Ryzen 9 3900X (Desktop) | 70 MB    | 12-24        | 3.4-4.6     |

Table: Some AMD processors

# MODERN SUPERCOMPUTERS

## What is a supercomputer?

- cdc 6600: 1964 – three million calculations per second
- Summit: 2018 – 36000 processors – 200 quadrillion calculations per second
- Frontier: 2022 – 8 million processors – AMD EPYC with 64 cores and speed up to 2GHz – quintillion calculations per second
- Toubkal: 2021 – 69000 processors



# MODERN SUPERCOMPUTERS

## What is a supercomputer?


Frontier (USA)



## MODERN SUPERCOMPUTERS

### What is a supercomputer?


Cluster



# MODERN SUPERCOMPUTERS


## Top 500

- Cray 2: Gigascale milestone in 1985
- Intel ASCI Red System: Terascale in 1997
- IBM Roadrunner System: Petascale in 2008
- Frontier: Exascale in 2022



# Modern supercomputers


## Top 500 Family system share evolution November 2009



1

# Modern supercomputers


## Top 500 Family system share evolution November 2011



1

# Modern supercomputers

## Top 500 Family system share evolution November 2015




1

# Modern supercomputers

## Top 500 Family system share evolution


November 2017



1

# Modern supercomputers

## Top 500 Family system share evolution June 2022



1



## SUMMARY

- Highlights
  - New architectures are available
  - Supercomputers achieve Exascale
- Consequence for the developers
  - Writing dedicated codes

## OUTLINE

- Some definitions
  - FLOPS
  - Frequency
  - Memory Bandwidth
  - Memory Latency
- Computational Intensity
- Two level memory model

## SOME DEFINITIONS

### FLOPS

Floating point operations per second (FLOPS or flop/second).

## SOME DEFINITIONS

### Frequency

Speed at which a processor or other component operates (Hz)

## SOME DEFINITIONS

### Memory Bandwidth

Rate at which data can be transferred between the CPU and the memory (bytes/second).

## SOME DEFINITIONS

### Memory Latency

Time delay between a processor requesting data from memory and the moment that the data is available for use (clock cycles or time units).

# COMPUTATIONAL INTENSITY

Algorithms have two costs (measured in time or energy):


- Arithmetic (FLOPS)
- Communication: moving data between
  - levels of a memory hierarchy (sequential case)
  - processors over a network (parallel case)

## Computational Intensity

It is the ratio between arithmetic complexity (or cost) and memory complexity (cost).

# TWO LEVEL MEMORY MODEL

## Modern architecture (CPU)



### Typical sizes


- RAM ~ 4 GB – 128 GB even higher on servers
- L3 ~ 4 MB – 50 MB
- L2 ~ 256 KB – 8 MB
- ➡ Holds data that is likely to be accessed by the CPU
- L1 ~ 256 KB
- ➡ Instruction and Data cache

### Cache Hit or Miss

- **Cache Hit:** CPU is able to find the Data in L1/L2/L3
- **Cache Miss:** CPU is not able to find the Data in L1-L2-L3 and must retrieve it from RAM



## MATRIX MULTIPLICATION: THREE NESTED LOOP



```
1 for i in range(0, n):  
2     #read row i of A into fast memory  
3     for j in range(0, n):  
4         #read row C[i,j] into fast memory  
5         #read col j of B into fast memory  
6         for k in range(0, n):  
7             C[i,j] = C[i,j] + A[i,k]*B[k,j]  
8             #write C[i,j] back to slow memory
```

```
1 arithmetic cost      :: n**3*(ADD + MUL) = 2n**3 arithmetic operations  
2 memory cost         :: n**3*READ + n**2*READ + n**2*(READ + WRITE) = n**3 + 3n**2  
3 computational intensity :: 2n**3/(n**3 + 3n**2) ~ 2
```


## SUMMARY

- Running time of an algorithm is sum of 3 terms:
  - `N_flops * time_per_flop`
  - `N_words / bandwidth`
  - `N_messages * latency`
- ➡ Avoiding communication algorithms come with a significant speedup
- Some examples
  - Up to 12x faster for 2.5D matmul on 64K core IBM BG/P
  - Up to 3x faster for tensor contractions on 2K core Cray XE/6
  - Up to 6.2x faster for All-Pairs-Shortest-Path on 24K core Cray CE6

## OUTLINE

- Data Locality
  - The Penalty of Stride
  - High Dimensional Arrays
- Block Matrix Multiplication

# DATA LOCALITY



- Data locality is key for improving per-core performance,
- Memory hierarchy has 4 levels,
- Processor looks for needed data in memory hierarchy,
- Simple or complex manipulations can increase speedup,
- Blocking version of `mxm` can increase computational intensity.

## The Penalty of Stride > 1?

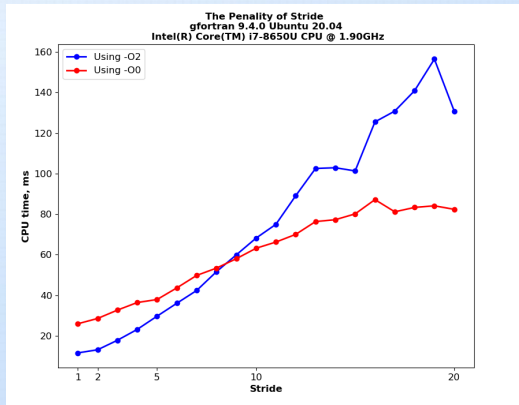
- Data should be arranged for unit stride access,
- Not doing so can result in severe performance penalty

### Example:

```
1 do i=1, N*i_stride,i_stride  
2     mean = mean + a(i)  
3 end do
```

- Compilation with all optimization and vectorization disabled (-O0)
- Compilation with (-O2) that activates some optimizations

## The Penalty of Stride: CPU time



## High Dimensional Arrays

|   |   |   |
|---|---|---|
| 1 | 2 | 3 |
| 4 | 5 | 6 |
| 7 | 8 | 9 |

Row-Major  
(C)

|  |  |   |   |   |   |   |   |   |   |   |  |
|--|--|---|---|---|---|---|---|---|---|---|--|
|  |  | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 |  |
|--|--|---|---|---|---|---|---|---|---|---|--|

Column-Major  
(Fortran)

|  |  |   |   |   |   |   |   |   |   |   |  |
|--|--|---|---|---|---|---|---|---|---|---|--|
|  |  | 1 | 4 | 7 | 2 | 5 | 8 | 3 | 6 | 9 |  |
|--|--|---|---|---|---|---|---|---|---|---|--|

- High Dimensional Arrays are stored as a contiguous sequence of elements
  - ➡ Fortran uses Column-Major ordering
  - ➡ C uses Row-Major ordering
- mxm in Fortran  $N = 1000$
- Naive version: CPU-time 1660.6 (msec)
  - Transpose version: CPU-time 1139.8 (msec)

# BLOCK MATRIX MULTIPLICATION

## mxm example: Using block version (cache optimization)


```
for ii in range(0, n, nb):
    for jj in range(0, n, nb):
        for kk in range(0, n, nb):
            for i in range(ii, min(ii+nb, n)):
                for j in range(jj, min(jj+nb, n)):
                    for k in range(kk, min(kk+nb, n)):
                        c[i][j] += a[i][k] * b[k][j]
```



**nb x nb**



**nb x nb**

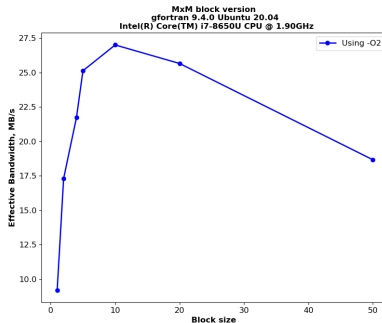
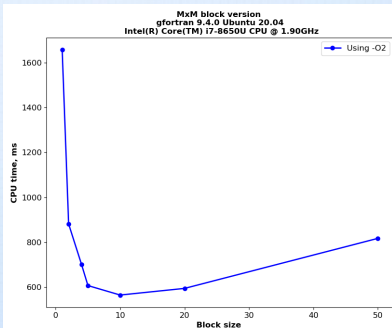


**nb x nb**



# BLOCK MATRIX MULTIPLICATION

## mxm block version: CPU time & Bandwidth



## SUMMARY

- Access contiguous, stride-one memory addresses
- Emphasize cache reuse
- Use data structures that improve locality
- Minimize communication across different memory levels
- Use parallelism to improve locality

## OUTLINE

- About Python
- Python is slow !
- Profiling a Python code

# ABOUT PYTHON

- Python was created by Guido van Rossum in 1991 (last version 3.11 – 24/10/2022)
- Python is simple
- Python is fully featured
- Python is readable
- Python is extensible
- Python is ubiquitous, portable, and free
- Python has many third party libraries, tools, and a large community

# ABOUT PYTHON

- Python was created by Guido van Rossum in 1991 (last version 3.11 – 24/10/2022)
- Python is simple
- Python is fully featured
- Python is readable
- Python is extensible
- Python is ubiquitous, portable, and free
- Python has many third party libraries, tools, and a large community

➡ But Python is slow!!

# ABOUT PYTHON

- Python was created by Guido van Rossum in 1991 (last version 3.11 – 24/10/2022)
- Python is simple
- Python is fully featured
- Python is readable
- Python is extensible
- Python is ubiquitous, portable, and free
- Python has many third party libraries, tools, and a large community

➡ When does it really matter?

# PYTHON IS SLOW

## When does it matter?

- Is my code fast?
- How many CPUh?
- Problems on the system?
- How much effort is it to make it run faster?

# 🔵 PROFILING A PYTHON CODE: WHY?

- Code bottlenecks
- Premature optimization is the root of all evil D. Knuth
- First make it work. Then make it right. Then make it fast. K. Beck
- How?



## ❏ PROFILING A PYTHON CODE: PROFILERS

- Deterministic and statistical profiling
  - the profiler will be monitoring all the events
  - it will sample after time intervals to collect that information
- The level at which resources are measured; module, function or line level
- Profile viewers

## PROFILING A PYTHON CODE: TOOLS

- Inbuilt timing modules
- profile and cProfile
- pstats
- line\_profiler
- snakeviz

## PROFILING A PYTHON CODE: USE CASE

```
1 def linspace(start, stop, n):
2     step =float(stop -start) / (n -1)
3     return [start +i *step for i in range(n)]
4
5 def mandel(c, maxiter):
6     z = c
7     for n in range(maxiter):
8         if abs(z) >2:
9             return n
10        z = z*z +c
11    return n
12
13 def mandel_set(xmin=-2.0, xmax=0.5, ymin=-1.25, ymax=1.25,
14               width=1000, height=1000, maxiter=80):
15     r = linspace(xmin, xmax, width)
16     i = linspace(ymin, ymax, height)
17     n = [[0]*width for _ in range(height)]
18     for x in range(width):
19         for y in range(height):
20             n[y][x] =mandel(complex(r[x], i[y]), maxiter)
21     return n
```

# PROFILING A PYTHON CODE: TIMEIT

The very naive way

```
1 import timeit
2
3 start_time =timeit.default_timer()
4 mandel_set()
5 end_time =timeit.default_timer()
6 # Time taken in seconds
7 elapsed_time =end_time -start_time
8
9 print('> Elapsed time', elapsed_time)
```

or using the magic method timeit

```
1 [In] %timeit mandel_set()
2 [Out] 3.01 s +/- 84.6 ms per loop (mean +/- std. dev. of 7 runs, 1 loop each)
```

# PROFILING A PYTHON CODE: PRUN

```
[In] %prun -s cumulative mandel_set()
```

which is, in console mode, equivalent to

```
python -m cProfile -s cumulative mandel.py
```

```
25214601 function calls in 5.151 seconds
```

```
Ordered by: cumulative time
```

| ncalls   | totttime | percall | cumtime | percall | filename:lineno(function)                        |
|----------|----------|---------|---------|---------|--|
| 1        | 0.000    | 0.000   | 5.151   | 5.151   | {built-in method builtins.exec}                  |
| 1        | 0.002    | 0.002   | 5.151   | 5.151   | <string>:1(<module>)                             |
| 1        | 0.291    | 0.291   | 5.149   | 5.149   | <ipython-input-4-9421bc2016cb>:13(mandel_set)    |
| 1000000  | 3.461    | 0.000   | 4.849   | 0.000   | <ipython-input-4-9421bc2016cb>:5(mandel)         |
| 24214592 | 1.388    | 0.000   | 1.388   | 0.000   | {built-in method builtins.abs}                   |
| 1        | 0.008    | 0.008   | 0.008   | 0.008   | <ipython-input-4-9421bc2016cb>:17(<listcomp>)    |
| 2        | 0.000    | 0.000   | 0.000   | 0.000   | <ipython-input-4-9421bc2016cb>:1(linspace)       |
| 2        | 0.000    | 0.000   | 0.000   | 0.000   | <ipython-input-4-9421bc2016cb>:3(<listcomp>)     |
| 1        | 0.000    | 0.000   | 0.000   | 0.000   | {method 'disable' of '_lsprof.Profiler' objects} |

# PROFILING A PYTHON CODE: LINE LEVEL

- Use the line\_profiler package

```
[In] %load_ext line_profiler
[In] %lprun -f mandel mandel_set()
```

```
Timer unit: 1e-06 s
Total time: 12.4456 s
File: <ipython-input-2-9421bc2016cb>
Function: mandel at line 5
#Line      Hits          Time  Per Hit   % Time  Line Contents
=====
5          5                def mandel(c, maxiter):
6      1000000      250304.0      0.3      1.1      z = c
7     24463110     6337732.0      0.3     27.7      for n in range(maxiter):
8     24214592     8327289.0      0.3     36.5          if abs(z) > 2:
9         751482      201108.0      0.3      0.9              return n
10     23463110     7658255.0      0.3     33.5          z = z*z + c
11     248518        65444.0      0.3      0.3      return n
```

## PROFILING A PYTHON CODE: LINE LEVEL

This can be done in console mode as well

```
@profile
def mandel(c, maxiter):
    z = c
    for n in range(maxiter):
        if abs(z) > 2:
            return n
        z = z*z + c
    return n
```

Then on the command line

```
kernprof -l -v mandel.py
```

Then

```
python3 -m line_profiler mandel.py.lprof
```

# PROFILING A PYTHON CODE: MEMORY

- Use the `memory_profiler` package

```
[In] %load_ext memory_profiler
[In] %mprun -f mandel mandel_set()
```

| Line # | Mem usage | Increment     | Occurrences | Line Contents            |
|--------|-----------|---------------|-------------|--------------------------|
| 8      | 118.2 MiB | -39057.7 MiB  | 1000000     | def mandel(c, maxiter):  |
| 9      | 118.2 MiB | -39175.5 MiB  | 1000000     | z = c                    |
| 10     | 118.2 MiB | -293081.8 MiB | 24463110    | for n in range(maxiter): |
| 11     | 118.2 MiB | -292425.7 MiB | 24214592    | if abs(z) > 2:           |
| 12     | 118.2 MiB | -38519.6 MiB  | 751482      | return n                 |
| 13     | 118.2 MiB | -253906.1 MiB | 23463110    | z = z*z + c              |
| 14     | 118.2 MiB | -656.4 MiB    | 248518      | return n                 |



## PROFILING A PYTHON CODE: MEMORY

- Use the `memory_profiler` package

```
@profile
def mandel(c, maxiter):
    z = c
    for n in range(maxiter):
        if abs(z) > 2:
            return n
        z = z*z + c
    return n
```

Then on the command line

```
mprof run mandel.py
```

Then

```
mprof plot
```

Or

```
python3 -m memory_profiler mandel.py
```

## OUTLINE

- Accelerate a Python code
  - Using Numpy
  - Using Cython
  - Using Numba
  - Using Pyccel
- Some Benchmarks

## ACCELERATE A PYTHON CODE: NUMPY


- Library for scientific computing in Python,
- High-performance multidimensional array object,
- Integrates C, C++, and Fortran codes in Python,
- Uses multithreading.

# ACCELERATE A PYTHON CODE: NUMPY VS LISTS

```
1 import numpy, time
2
3 size =1000000
4
5 print("Concatenation: ")
6 list1 =[i for i in range(size)]; list2 =[i for i in range(size)]
7
8 array1 =numpy.arange(size); array2 =numpy.arange(size)
9
10 # List
11 initialTime =time.time()
12 list1 =list1 +list2
13 # calculating execution time
14 print("Time taken by Lists :", (time.time() -initialTime), "seconds")
15
16 # Numpy array
17 initialTime =time.time()
18 array =numpy.concatenate((array1, array2), axis =0)
19 # calculating execution time
20 print("Time taken by NumPy Arrays :", (time.time() -initialTime), "seconds")
```

```
1 Concatenation:
2 Time taken by Lists : 0.021048307418823242 seconds
3 Time taken by NumPy Arrays : 0.009451150894165039 seconds
```

# ACCELERATE A PYTHON CODE: CYTHON



- Cython is an optimizing static compiler for:
  - Python programming language
  - Cython programming language (based on Pyrex)
- Cython gives you the combined power of Python.

# ACCELERATE A PYTHON CODE: CYTHON

- Python

```
1 def mandelbrot(m, size, iterations):  
2     for i in range(size):  
3         for j in range(size):  
4             c = -2 + 3./size*j + 1j*(1.5-3./size*i)  
5             z = 0  
6             for n in range(iterations):  
7                 if np.abs(z) <=10:  
8                     z = z*z + c; m[i, j] = n  
9                 else:  
10                    break
```

# ACCELERATE A PYTHON CODE: CYTHON

- Cython

```
1 def mandelbrot_cython(int[:,::1] m,int size, int iterations):
2     cdef int i, j, n
3     cdef complex z, c
4     for i in range(size):
5         for j in range(size):
6             c = -2 + 3./size*j + 1j*(1.5-3./size*i)
7             z = 0
8             for n in range(iterations):
9                 if z.real**2 + z.imag**2 <= 100:
10                     z = z*z + c; m[i, j] = n
11             else:
12                 break
```

# ACCELERATE A PYTHON CODE: CYTHON

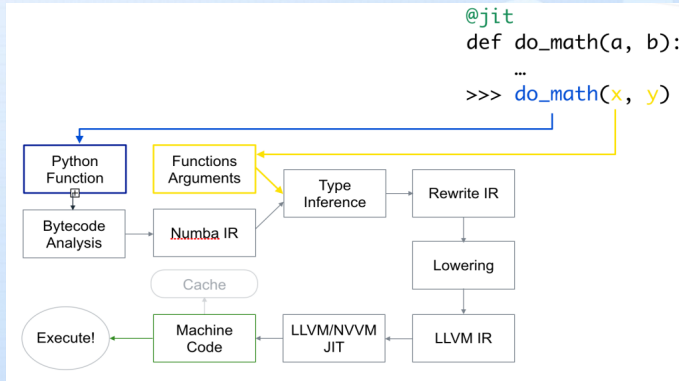
- Execution time

```
1 %%timeit -n1 -r1
2 m = np.zeros(s, dtype=np.int32)
3 mandelbrot(m, size, iterations)
4 >> 12.2 s +/- 0 ns per loop (mean +/- std. dev. of 1 run, 1 loop each)
5
6
7 %%timeit -n1 -r1
8 m = np.zeros(s, dtype=np.int32)
9 mandelbrot_cython(m, size, iterations)
10 >> 29.8 ms +/- 0 ns per loop (mean +/- std. dev. of 1 run, 1 loop each)
```



# ACCELERATE A PYTHON CODE: NUMBA

- Open source Just-In-Time compiler for python functions.
- Uses the LLVM library as the compiler backend.



# ACCELERATE A PYTHON CODE: NUMBA

- Python

```
1 import numpy as np
2
3 def do_sum():
4     acc = 0.
5     for i in range(10000000) :
6         acc += np.sqrt(i)
7     return acc
```

- Numba

```
1 from numba import njit
2
3 @njit
4 def do_sum_numba():
5     acc = 0.
6     for i in range(10000000) :
7         acc += np.sqrt(i)
8     return acc
```

```
1 Time for Pure Python Function: 7.724030017852783
```

```
2 Time for Numba Function: 0.015453100204467773
```

## ACCELERATE A PYTHON CODE: PYCCCEL

- Pyccel is a static compiler for Python 3, using Fortran or C as a backend language.
- Python function:

```
1 import numpy as np
2
3 def do_sum_pyccel():
4     acc = 0.
5     for i in range(10000000) :
6         acc += np.sqrt(i)
7     return acc
```

# ACCELERATE A PYTHON CODE: PYCCEL (F90)

- Compilation using fortran:

```
pyccel --language=fortran pyccel_example.py
```

```
1 module pyccel_example
2 use, intrinsic :: ISO_C_Binding, only : i64 => C_INT64_T , f64 => C_DOUBLE
3 implicit none
4 contains
5 !.....
6 function do_sum_pyccel() result(acc)
7
8     implicit none
9     real(f64) :: acc
10    integer(i64) :: i
11    acc = 0.0_f64
12    do i = 0_i64, 9999999_i64, 1_i64
13        acc = acc + sqrt(Real(i, f64))
14    end do
15    return
16 end function do_sum_pyccel
17 !.....
18 end module pyccel_example
```

```
1 Time for Pure Python Function: 7.400242328643799
```

```
2 Time for Pyccel Function: 0.01545262336730957
```

# ACCELERATE A PYTHON CODE: PYCCEL (C)

- Compilation using c:

```
pyccel --language=c pyccel_example.py
```

```
1  #include "pyccel_example.h"
2  #include <stdlib.h>
3  #include <math.h>
4  #include <stdint.h>
5  /*.....*/
6  double do_sum_pyccel(void)
7  {
8      int64_t i;
9      double acc;
10     acc = 0.0;
11     for (i = 0; i < 100000000; i += 1)
12     {
13         acc += sqrt((double)(i));
14     }
15     return acc;
16 }
17 /*.....*/
```

## SOME BENCHMARKS

### Rosen-Der

| Tool              | Python | Cython          | Numba          | Pythran         | Pyccel-gcc      | Pyccel-intel    |
|-------------------|--------|-----------------|----------------|-----------------|-----------------|-----------------|
| Timing ( $\mu$ s) | 229.85 | 2.06            | 4.73           | 2.07            | 0.98            | 0.64            |
| Speedup           | —      | $\times 111.43$ | $\times 48.57$ | $\times 110.98$ | $\times 232.94$ | $\times 353.94$ |

### Black-Scholes

| Tool              | Python | Cython        | Numba          | Pythran        | Pyccel-gcc      | Pyccel-intel         |
|-------------------|--------|---------------|----------------|----------------|-----------------|----------------------|
| Timing ( $\mu$ s) | 180.44 | 309.67        | 3.0            | 1.1            | 1.04            | $6.56 \cdot 10^{-2}$ |
| Speedup           | —      | $\times 0.58$ | $\times 60.06$ | $\times 163.8$ | $\times 172.35$ | $\times 2748.71$     |

### Laplace

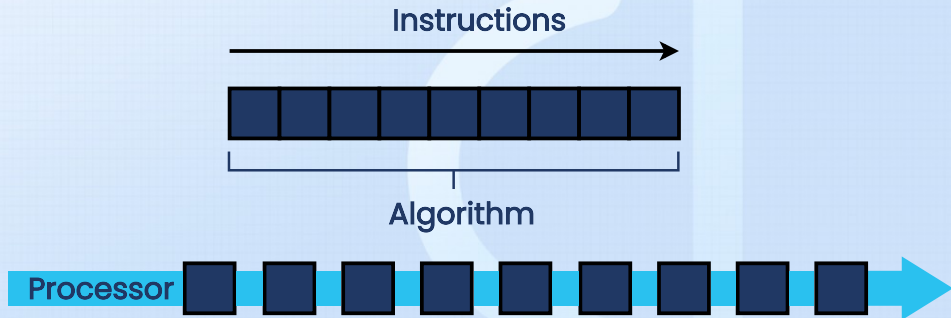
| Tool              | Python | Cython        | Numba                | Pythran              | Pyccel-gcc           | Pyccel-intel         |
|-------------------|--------|---------------|----------------------|----------------------|----------------------|----------------------|
| Timing ( $\mu$ s) | 57.71  | 7.98          | $6.46 \cdot 10^{-2}$ | $6.28 \cdot 10^{-2}$ | $8.02 \cdot 10^{-2}$ | $2.81 \cdot 10^{-2}$ |
| Speedup           | —      | $\times 7.22$ | $\times 892.02$      | $\times 918.56$      | $\times 719.32$      | $\times 2048.65$     |

## OUTLINE

- Parallel Programming
- Parallel strategies
- Parallel infrastructures

## PARALLEL PROGRAMMING

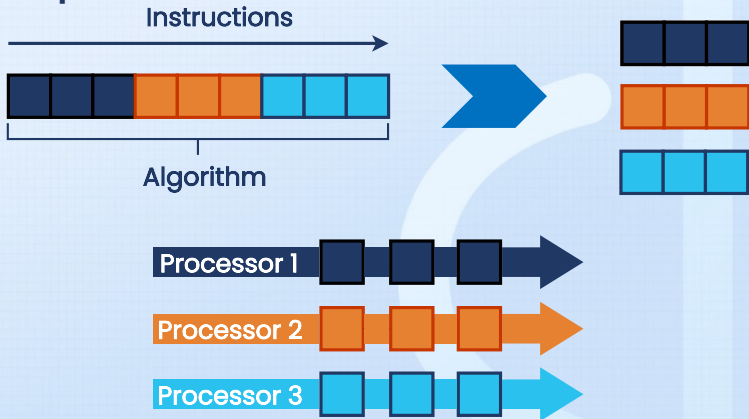
### Serial problem






## PARALLEL PROGRAMMING

### Parallel problem



## PARALLEL STRATEGIES

Setting



## PARALLEL STRATEGIES

### Setting



### Data parallelism




## PARALLEL STRATEGIES

### Setting




### Task parallelism




## PARALLEL INFRASTRUCTURES

### Shared memory



## PARALLEL INFRASTRUCTURES

### Distributed memory



## SUMMARY

- Advantages
  - Parallel computing saves time
  - Solve Larger Problems
  - Data storage
- Walls
  - Writing parallel codes is time-consuming
  - Power consumption
  - CPUs, GPUs, FPGAs

## OUTLINE


- Performance and scalability
- Classes of performance metrics
- Execution time & Total Parallel Overhead
- Speedup & Efficiency
- Amdahl's & Gustafson's Laws



## PERFORMANCE AND SCALABILITY

Design of parallel applications:

- Performance
- Scalability



## CLASSES OF PERFORMANCE METRICS

Distinct classes of performance metrics:

- Performance metrics for processors/cores
- Performance metrics for parallel applications
  - Practical metrics
  - Theoretical metrics

## ❏ EXECUTION TIME


**Finding minimum element among [23,12,9,30,37,27,8,17].**

- Execution time:
  - Serial Time  $T_s (\theta(n))$

## EXECUTION TIME

**Finding minimum element among [23,12,9,30,37,27,8,17].**

- Execution time:
  - Serial Time  $T_s (\theta(n))$
  - Parallel Time  $T_p (\theta(\log n))$



## ❑ TOTAL PARALLEL OVERHEAD

**Finding minimum element among [23,12,9,30,37,27,8,17].**

- Execution time:
  - Serial Time  $T_s (\theta(n))$
  - Parallel Time  $T_p (\theta(\log n))$

### **Total Parallel Overhead**

- With  $p$  processes:
  - Total time =  $pT_p$
  - Overhead =  $pT_p - T_s$

## SPEEDUP

Speedup is a measure of performance.

$$S_p = \frac{T_s}{T_p}$$

- Example 1: Find out the minimum element in array

$$S_p = \frac{\theta(n)}{\theta(\log n)} = \theta\left(\frac{n}{\log n}\right)$$

- Example 2: Solve 1D transport equation

|        | 1 CPUs | 2 CPUs | 4 CPUs | 8 CPUs | 16 CPUs |
|--------|--------|--------|--------|--------|---------|
| $T(p)$ | 1000   | 520    | 280    | 160    | 100     |
| $s(p)$ | 1      | 1.92   | 3.57   | 6.25   | 10.00   |

## EFFICIENCY

Efficiency is a measure of the usage of the computational capacity.

$$E_p = \frac{S_p}{p} = \frac{T_s}{p \times T_p}$$

- Example 1: Find out the minimum element in array

$$E_p = \frac{\theta\left(\frac{n}{\log n}\right)}{p} \quad (\text{if } p = n) \Rightarrow E_p = \frac{\theta\left(\frac{n}{\log n}\right)}{n} = \theta\left(\frac{1}{\log n}\right)$$

- Example 2: Solve 1D transport equation

|        | 1 CPUs | 2 CPUs | 4 CPUs | 8 CPUs | 16 CPUs |
|--------|--------|--------|--------|--------|---------|
| $S(p)$ | 1      | 1.92   | 3.57   | 6.25   | 10.00   |
| $E(p)$ | 1      | 0.96   | 0.89   | 0.78   | 0.63    |

## AMDAHL'S LAW: STRONG SCALING

- Serial part:  $0 \leq f \leq 1$

$$S_p = \frac{1}{f + \frac{1-f}{p}}$$

- Example: If  $f = 10\%$ , what is the max. speedup achievable using 8 processors?

- Solution:

$$S_p = \frac{1}{0.1 + \frac{1-0.1}{8}} \approx 4.7$$

- Limit:

$$\lim_{p \rightarrow \infty} = \frac{1}{0.1 + \frac{1-0.1}{p}} = 10$$



## ■ GUSTAFSON'S LAW: WEAK SCALING

- Parallel part:  $0 \leq f' \leq 1$

$$S_p = (1 - f') + f' \times p = 1 + (p - 1) \times f'$$

- Example:  $f' = 90\%$ , what is the scaled speedup using 64 processors?
- Solution:

$$S_p = 1 + (p - 1) \times f' = 1 + (64 - 1) \times 0.90 = 57.70$$

## SUMMARY

- Scalable application
  - Strong scaling + Weak scaling

|            |            | 1 CPUs | 2 CPUs | 4 CPUs | 8 CPUs | 16 CPUs |
|------------|------------|--------|--------|--------|--------|---------|
| Efficiency | n = 10 000 | 1      | 0.81   | 0.53   | 0.28   | 0.16    |
|            | n = 20 000 | 1      | 0.94   | 0.80   | 0.59   | 0.42    |
|            | n = 40 000 | 1      | 0.96   | 0.89   | 0.79   | 0.58    |

- Superlinear Speedup

$$\frac{T_s}{T_p} \geq p$$